

FACULTY OF MATHEMATICS,
PHYSICS AND INFORMATICS
Comenius University
Bratislava

Neural Networks for Computer Vision

Lecture 2: Image Classification

Ing. Viktor Kocur, PhD., RNDr. Zuzana Černeková, PhD.

27.9.2022

Acknowledgment



The majority of slides are directly adopted from slides for CS231n¹ course at Stanford University!

¹Stanford CS231n lecture slides. <http://cs231n.stanford.edu/slides/>

Contents



- Image Classification Task
- Challenges
- kNN Classifier
- Hyperparameters
- Linear Classifier

Image Classification

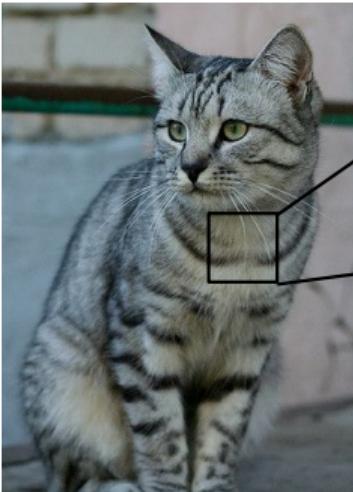


Image Classification



⇒ Cat

The Problem: Semantic Gap



This image by [Nikita](#) is
licensed under [CC-BY 2.0](#)

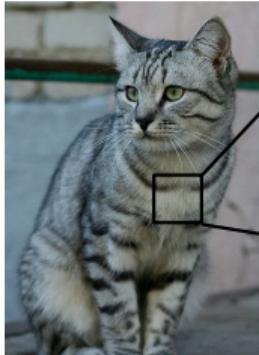
[185 112 188 111 104 99 106 99 96 103 112 119 104 97 93 87]
[91 98 182 186 104 79 98 183 99 185 123 136 119 185 94 85]
[76 85 90 185 128 185 87 96 95 99 115 112 106 103 99 85]
[99 81 81 93 128 131 127 108 95 98 182 99 96 93 101 94]
[106 91 61 64 69 91 88 85 181 107 109 98 75 84 96 95]
[114 108 85 55 59 69 64 60 54 87 112 125 98 74 84 87]
[114 130 147 183 65 81 88 65 52 80 77 84 182 93 85 82]
[128 137 144 180 69 85 86 66 52 80 77 84 182 93 85 82]
[125 133 148 137 119 121 117 79 85 79 80 65 54 64 72 88]
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
[115 114 109 123 158 148 131 118 113 109 109 92 74 65 72 78]
[89 93 96 97 100 147 131 118 113 114 113 103 106 95 77 88]
[63 77 80 85 77 79 102 123 117 115 117 123 125 138 115 87]
[62 65 82 89 70 71 88 181 124 126 119 181 107 114 131 119]
[63 65 75 88 89 71 62 81 128 138 135 105 81 98 118 118]
[87 65 71 87 106 95 69 45 76 138 126 107 92 94 105 112]
[118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
[164 146 112 88 82 128 124 184 76 48 45 66 88 101 102 109]
[157 170 157 128 93 86 114 132 112 97 69 55 78 82 99 94]
[130 128 134 161 139 188 109 118 121 134 114 87 65 53 69 86]
[128 112 96 117 158 144 128 115 184 107 102 93 87 81 72 79]
[123 107 96 86 83 112 153 149 122 109 104 75 88 107 112 99]
[122 121 182 80 82 86 94 117 145 148 153 102 58 78 92 107]
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84])]

What the computer sees

An image is a tensor of integers
between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Challenges: Viewpoint variation

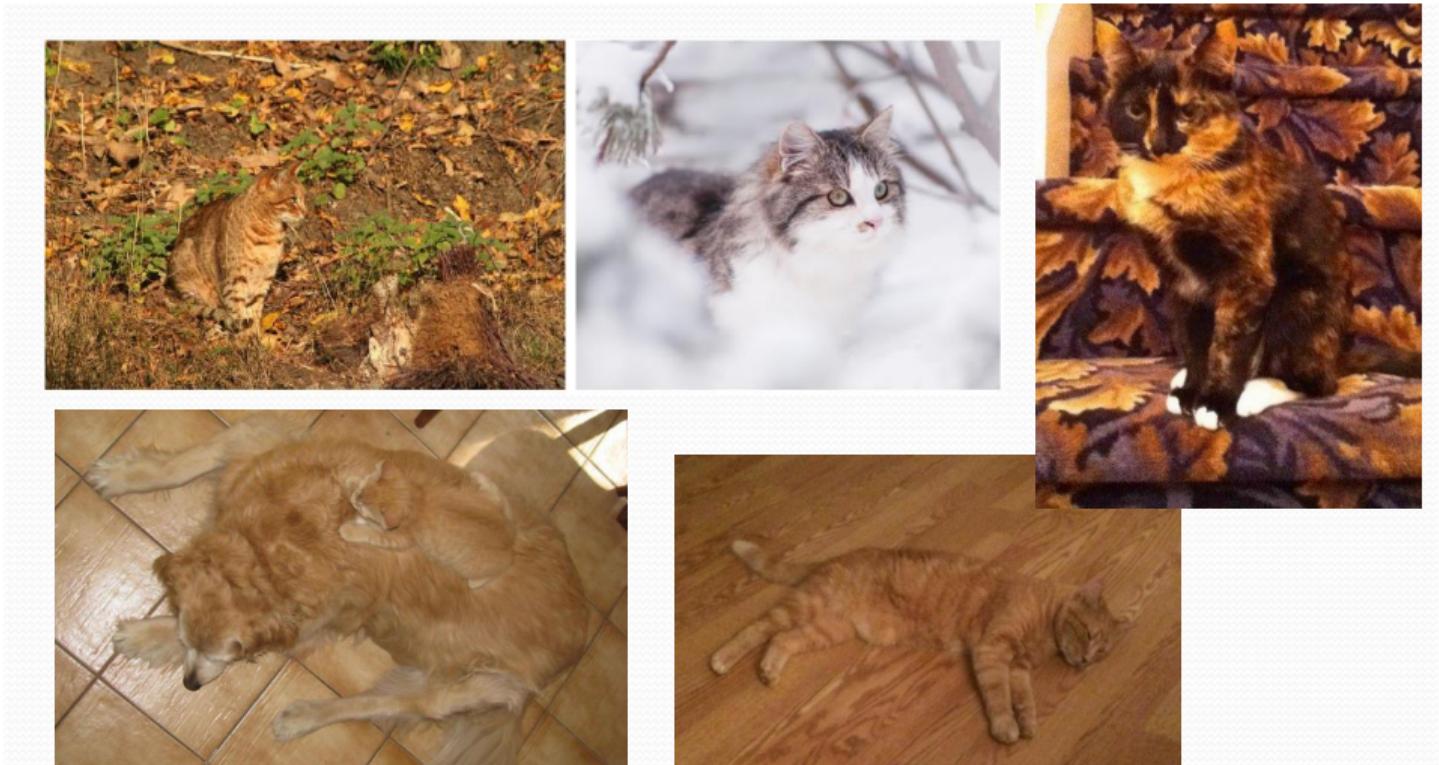


1185	112	183	111	184	99	186	99	99	183	112	118	184	97	93	971
1186	112	183	111	184	99	186	99	99	183	112	118	184	97	93	971
1186	95	99	185	128	186	97	182	99	99	115	112	186	180	99	951
1199	91	91	93	128	181	127	189	99	99	182	99	99	981	941	
1199	91	91	93	128	181	127	189	99	99	182	99	99	981	941	
1199	95	95	95	95	95	95	95	95	95	95	95	95	95	95	951
1114	189	95	95	95	95	95	95	95	95	95	95	95	95	95	951
1135	137	147	183	85	81	88	65	52	34	74	84	182	95	89	821
1129	125	137	137	137	137	137	137	94	95	79	88	85	64	72	1011
1129	125	137	137	137	137	137	137	94	95	79	88	85	64	72	1011
1127	125	131	147	133	127	128	131	111	98	89	75	81	64	72	941
1115	113	129	137	137	137	137	137	137	111	111	111	111	111	111	111
1101	99	99	97	148	147	131	118	112	114	113	186	186	95	77	881
1101	99	99	97	148	147	131	118	112	117	117	117	117	115	115	871
1103	77	95	95	95	95	95	95	95	95	95	95	95	95	95	951
1103	63	65	75	88	89	73	62	61	129	139	135	185	81	98	1181
1103	87	85	75	87	184	95	69	49	76	139	128	187	92	94	1181
1118	112	112	99	99	99	99	99	99	99	99	99	99	99	99	991
1164	146	112	99	99	82	239	124	186	79	98	45	86	86	182	1891
1157	178	157	128	93	88	113	132	112	97	99	55	79	82	99	941
1120	112	99	117	158	144	128	115	184	187	187	92	87	81	72	791
1120	112	99	117	158	144	128	115	184	187	187	92	87	81	72	791
1123	187	95	88	83	112	153	149	122	189	184	75	88	107	112	991
1123	187	95	88	83	112	153	149	122	189	184	75	88	107	112	991
1122	164	149	183	75	56	78	83	93	183	118	139	182	61	69	1071
1122	164	149	183	75	56	78	83	93	183	118	139	182	61	69	1071

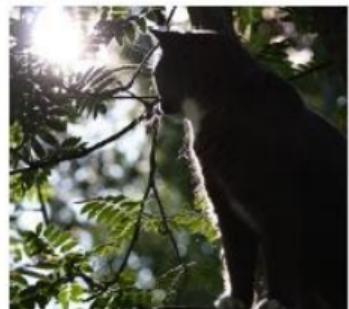


All pixels change when
the camera moves!

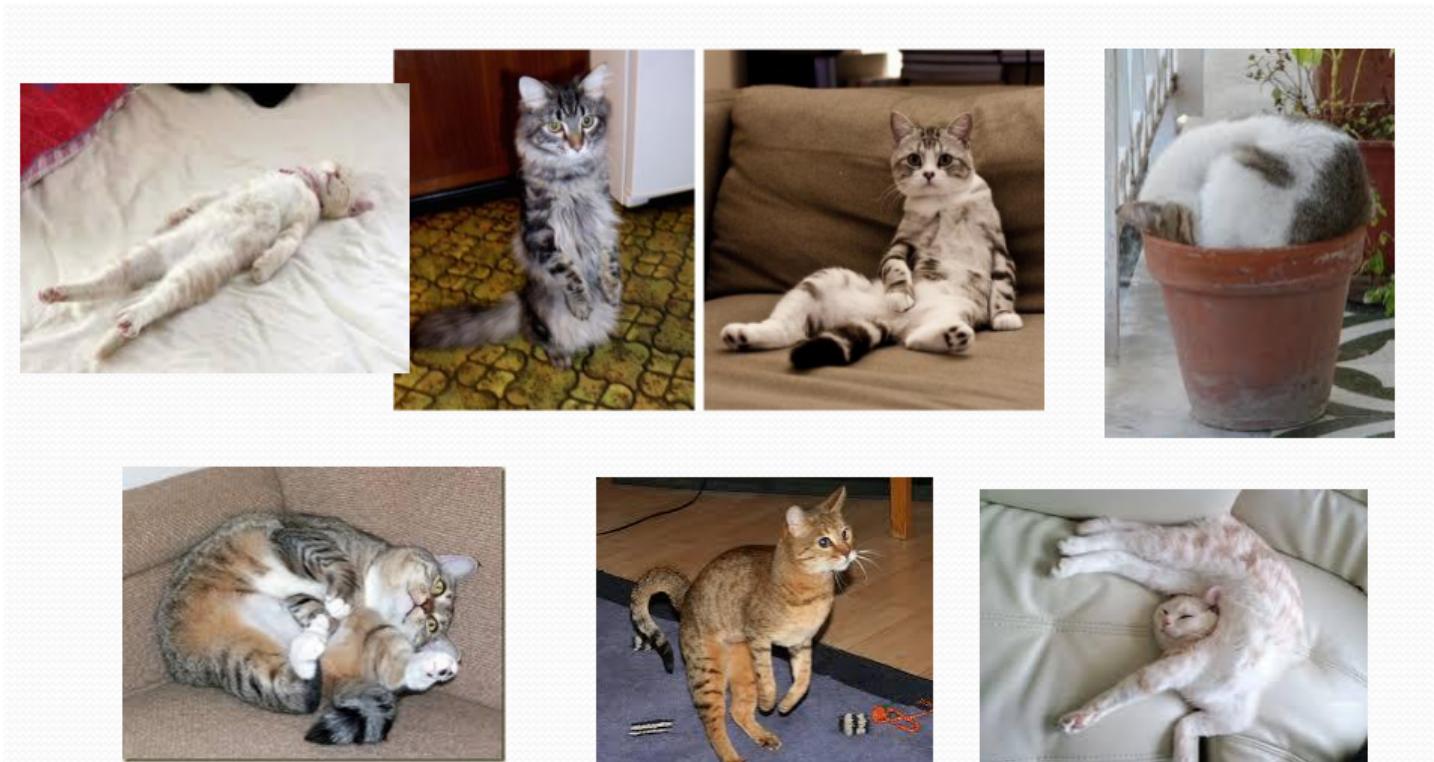
Challenges - background



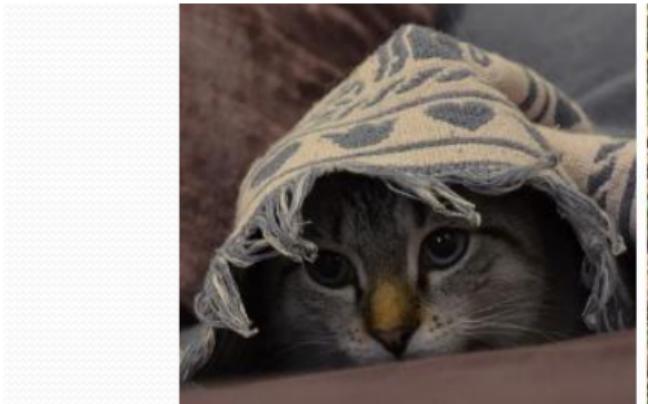
Challenges - illumination



Challenges - deformation



Challenges - occlusion



Challenges - intraclass variation



Image Classifier



```
def classify_image(image):
    # some magic
    return class_label
```

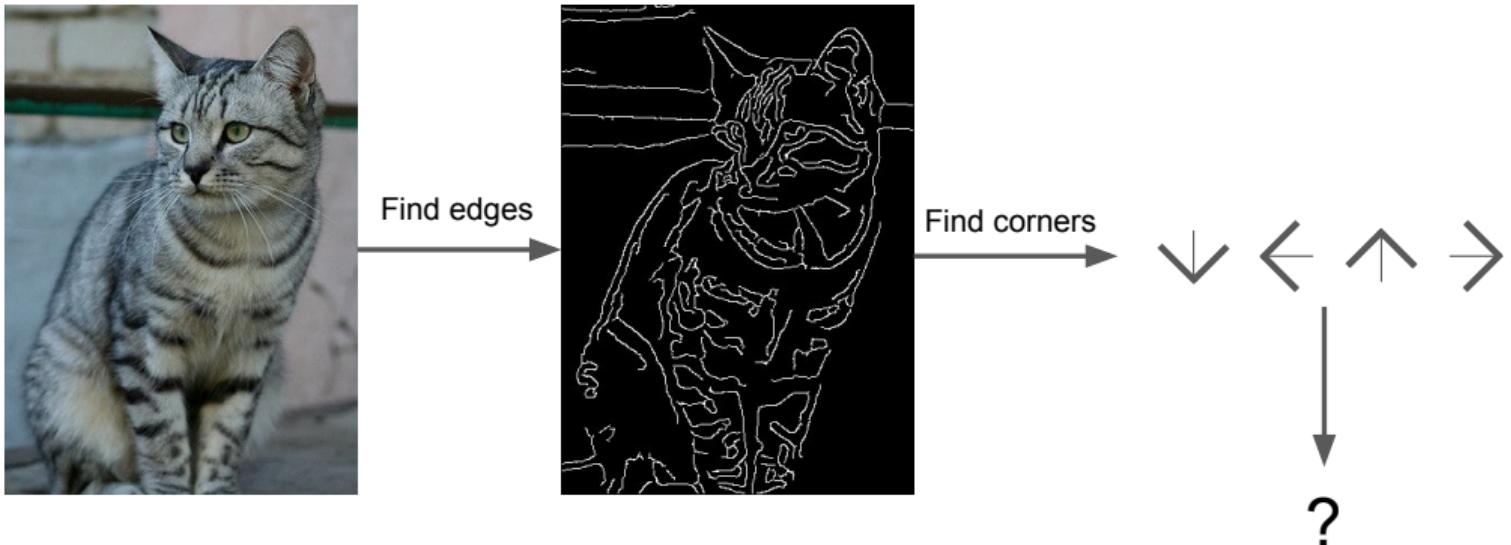
Image Classifier



```
def classify_image(image):  
    # some magic  
    return class_label
```

There is no obvious way how to do this!

Classical CV



John Canny, "A Computational Approach to Edge Detection", IEEE TPAMI 1986

Machine Learning



Machine learning algorithms are algorithms which can improve their performance from experience. In case of CV experience usually means large annotated datasets.

Typical ML workflow:

1. Collect a dataset of images and corresponding labels
2. Train a classifier using ML
3. Evaluate the classifier on a new set of images

Nearest Neighbors classifier



One of the simplest classifiers is the nearest neighbor classifier:

- We train the classifier by simply remembering all of the training images and labels.
- To classify a new image we assign it the label of the image from the training set which is closest the new image.

Distance metric



Training data with labels



query data

Distance Metric |  | $\rightarrow \mathbb{R}$

Simple metric

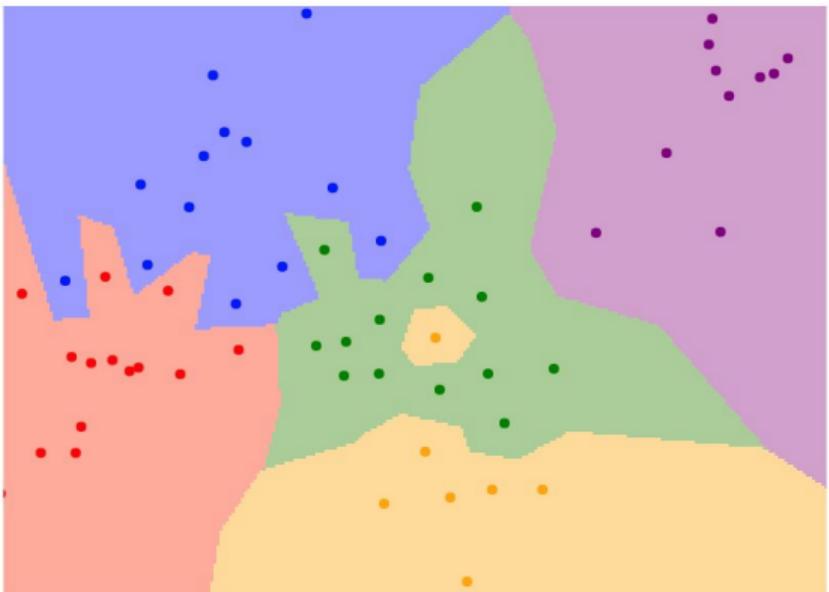


L1 distance: $d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$

test image				training image				pixel-wise absolute value differences			
56	32	10	18	10	20	24	17	46	12	14	1
90	23	128	133	8	10	89	100	82	13	39	33
24	26	178	200	12	16	178	170	12	10	0	30
2	0	255	220	4	32	233	112	2	32	22	108

- = add → 456

Decision boundary



1-nearest neighbor

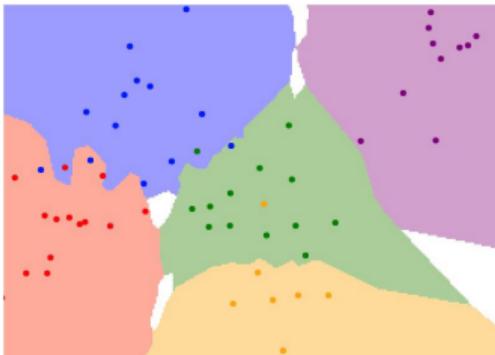
k-Nearest Neighbors



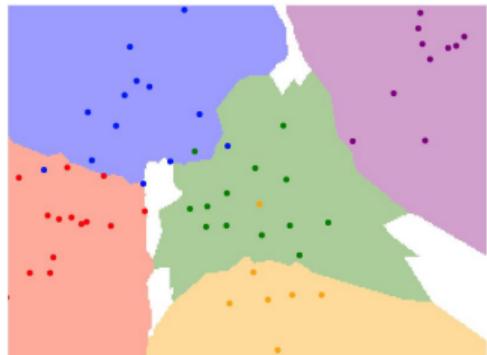
Instead of copying label from nearest neighbor,
take **majority vote** from K closest points



K = 1



K = 3



K = 5

Hyperparameters



Before using kNN we have to make few decisions:

- Which k do we use? E.g. number of neighbors.
- What kind of metric we use to calculate the distances?

Hyperparameters



Before using kNN we have to make few decisions:

- Which k do we use? E.g. number of neighbors.
- What kind of metric we use to calculate the distances?

Both of these are **hyperparameters**. In general, hyperparameters change the algorithms themselves. We use the prefix hyper, since more sophisticated algorithms usually have parameters which are learned during training.

Setting these is difficult and dependent on the given problem! Usually you have to try and see what works best!

Setting Hyperparameters



Idea #1: Choose hyperparameters
that work best on the **training data**

train

Setting Hyperparameters



Idea #1: Choose hyperparameters
that work best on the **training data**

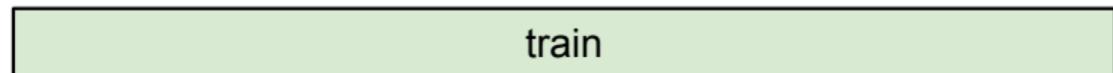
BAD: $K = 1$ always works
perfectly on training data

train

Setting Hyperparameters



Idea #1: Choose hyperparameters that work best on the **training data**



BAD: $K = 1$ always works perfectly on training data

Idea #2: choose hyperparameters that work best on **test data**



Setting Hyperparameters



Idea #1: Choose hyperparameters that work best on the **training data**

BAD: $K = 1$ always works perfectly on training data

train

Idea #2: choose hyperparameters that work best on **test data**

BAD: No idea how algorithm will perform on new data

train

test

Never do this!

Setting Hyperparameters



Idea #1: Choose hyperparameters that work best on the **training data**

BAD: $K = 1$ always works perfectly on training data

train

Idea #2: choose hyperparameters that work best on **test data**

BAD: No idea how algorithm will perform on new data

train

test

Idea #3: Split data into **train**, **val**; choose hyperparameters on val and evaluate on test

Better!

train

validation

test

Setting Hyperparameters



train

Idea #4: Cross-Validation: Split data into **folds**,
try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

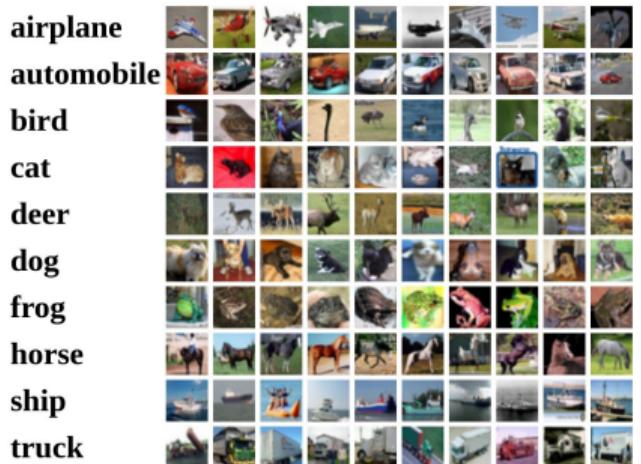
Example Dataset - CIFAR10



10 classes

50,000 training images

10,000 testing images



Alex Krizhevsky, "Learning Multiple Layers of Features from Tiny Images", Technical Report, 2009.



Example Dataset - CIFAR10

10 classes

50,000 training images

10,000 testing images

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



Test images and nearest neighbors



Alex Krizhevsky, "Learning Multiple Layers of Features from Tiny Images", Technical Report, 2009.



k-Nearest Neighbor with pixel distance **never used**.

- Distance metrics on pixels are not informative
- Very slow at test time



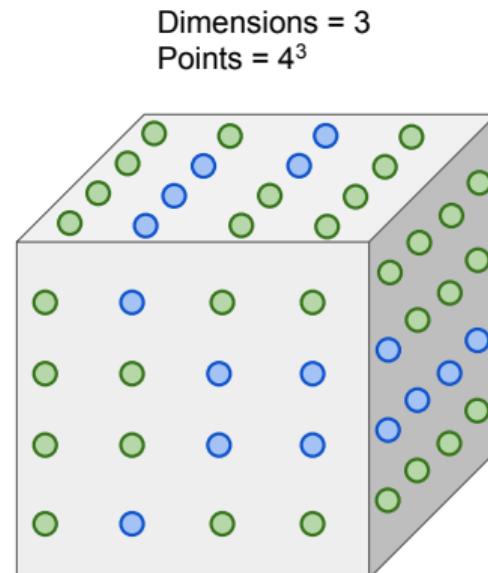
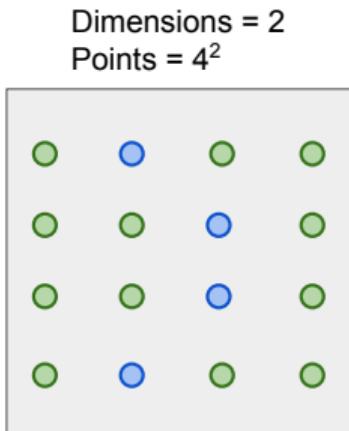
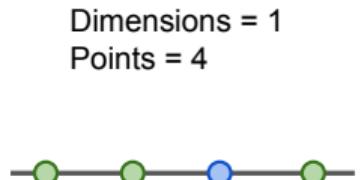
Original image is
CC0 public domain

kNN - not a good choice



k-Nearest Neighbor with pixel distance **never used**.

- Curse of dimensionality



Linear Classifier



Array of **32x32x3** numbers
(3072 numbers total)

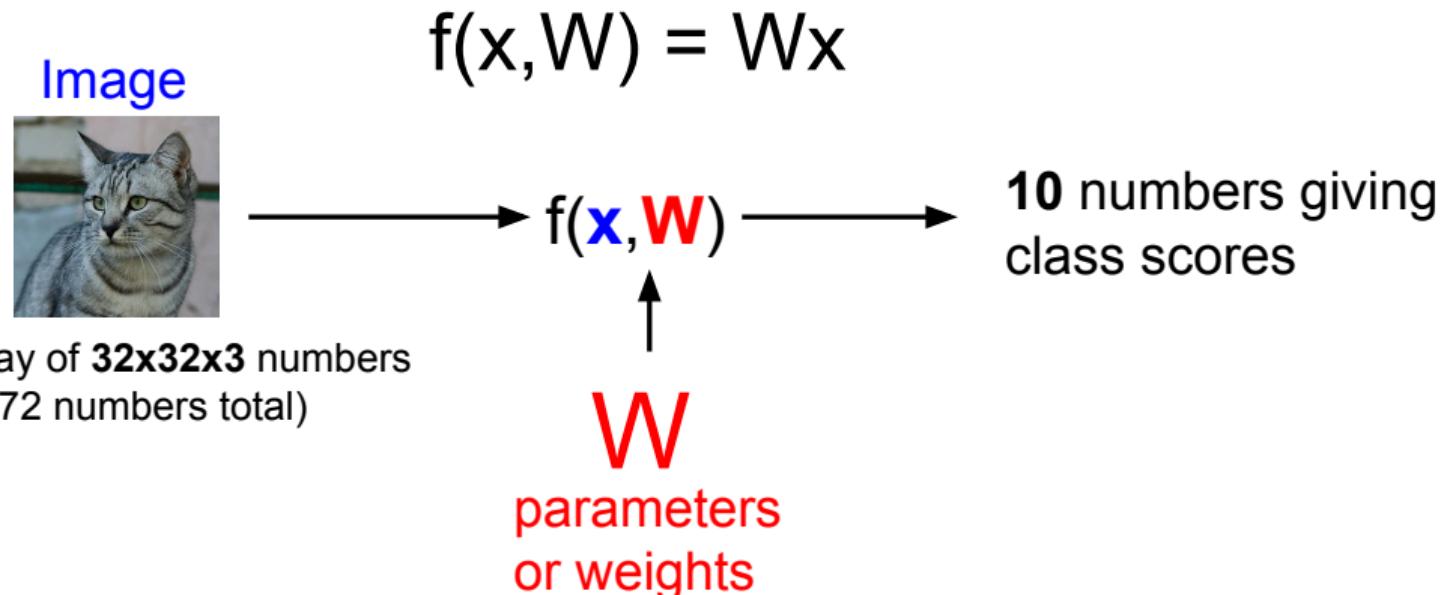
$$\rightarrow f(\mathbf{x}, \mathbf{W}) \rightarrow$$



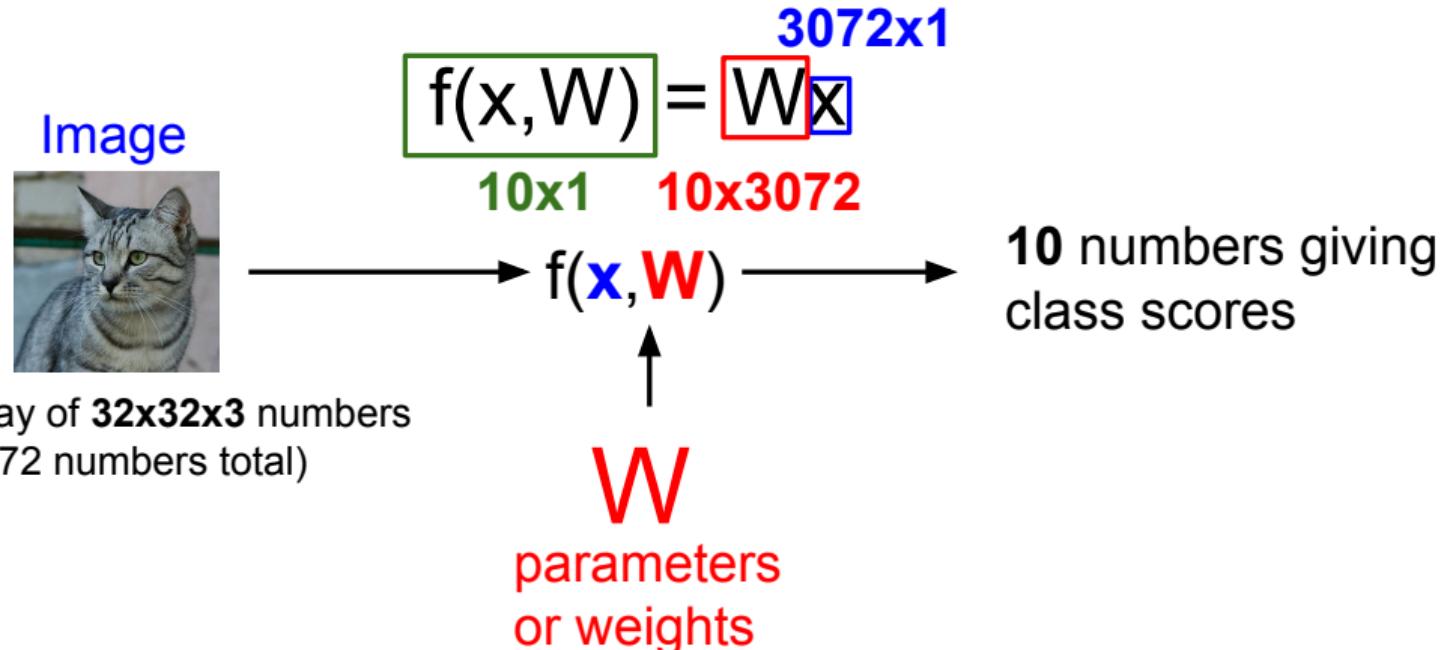
W
parameters
or weights

10 numbers giving
class scores

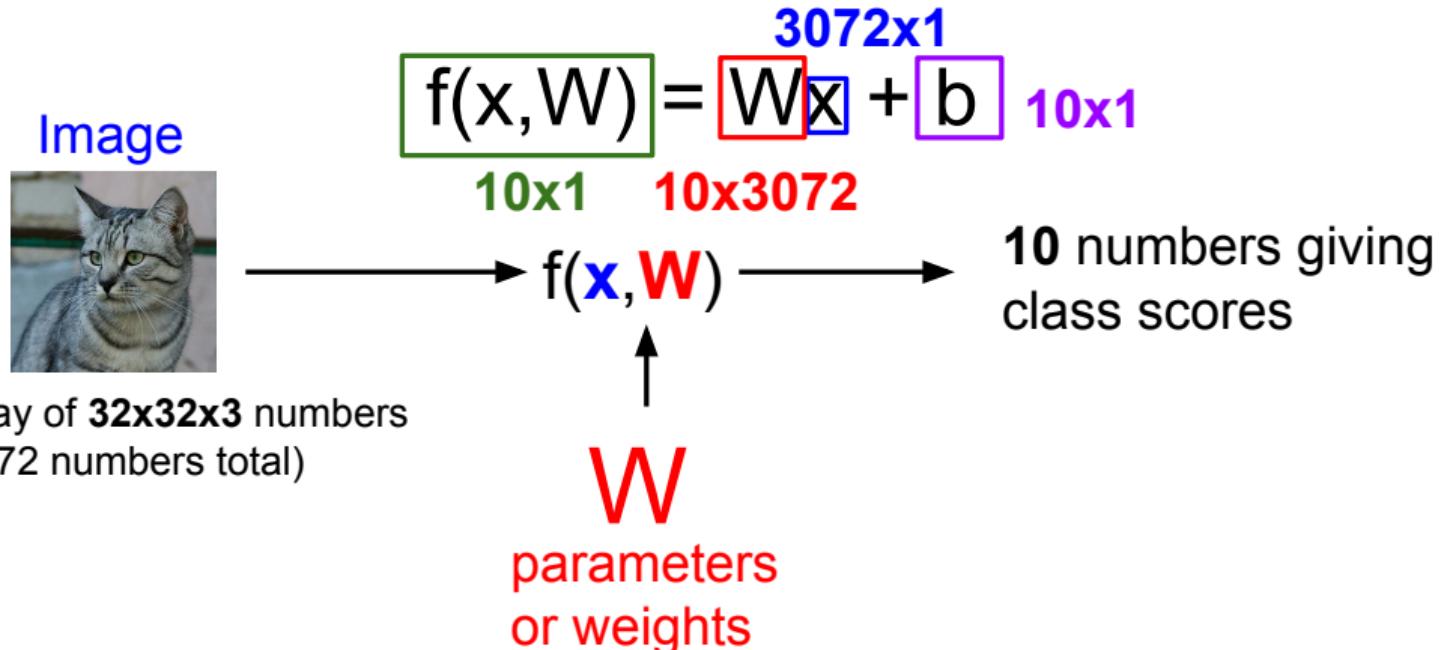
Linear Classifier



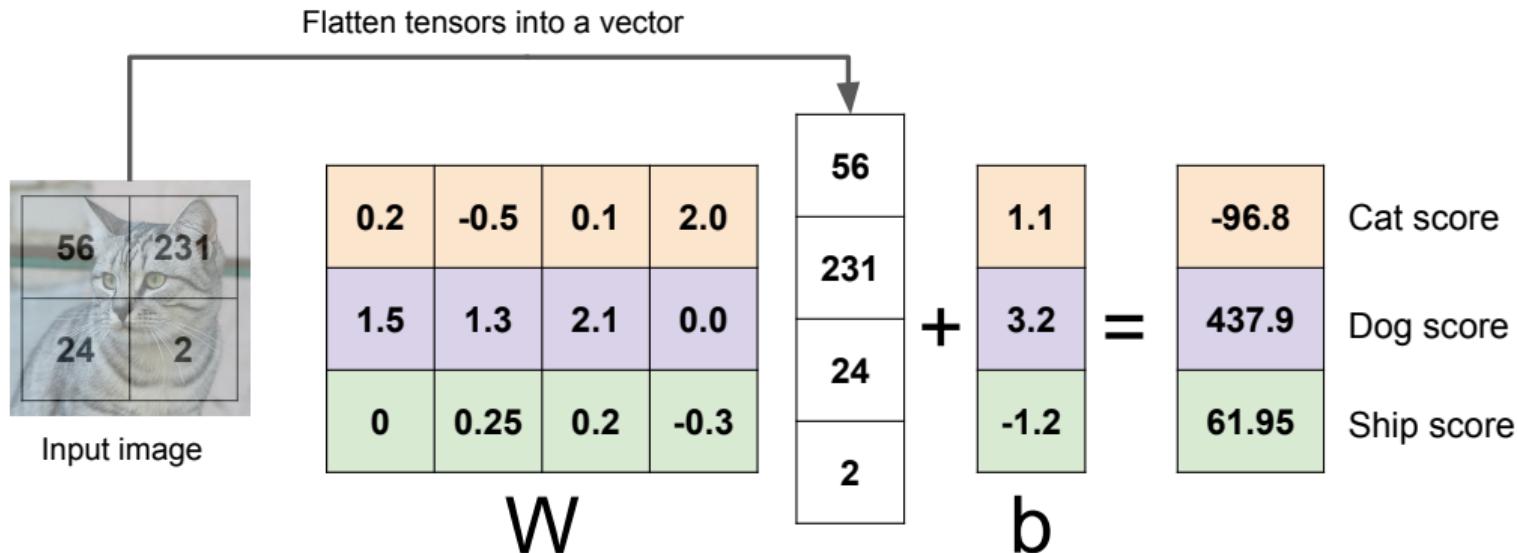
Linear Classifier



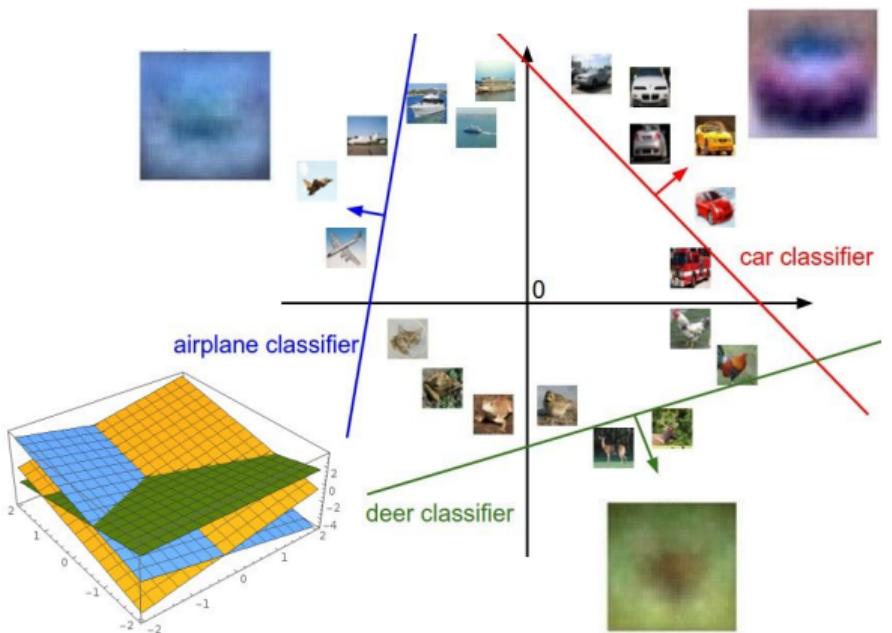
Linear Classifier



Linear Classifier - Example



Linear Classifier - Decision Boundary



$$f(x, W) = Wx + b$$



Array of **32x32x3** numbers
(3072 numbers total)

Linear Classifier - Hard cases

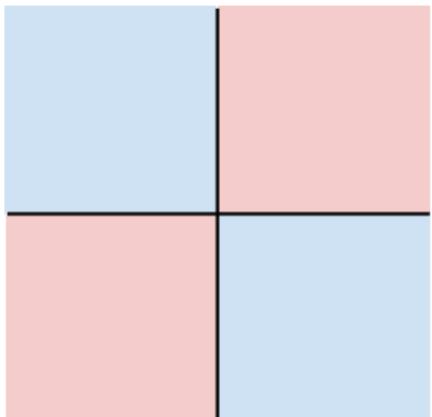


Class 1:

First and third quadrants

Class 2:

Second and fourth quadrants

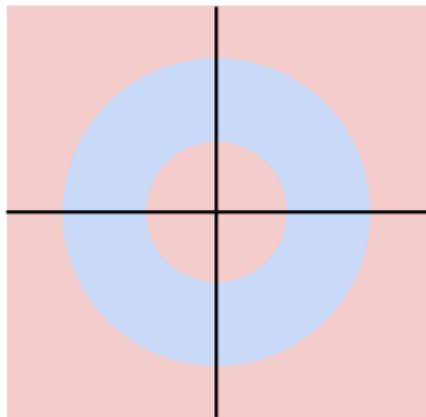


Class 1:

$1 \leq L_2 \text{ norm} \leq 2$

Class 2:

Everything else

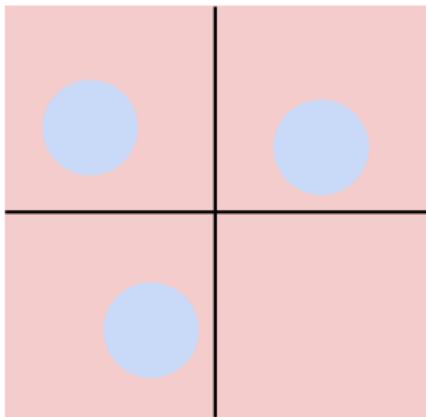


Class 1:

Three modes

Class 2:

Everything else





$$f(\mathbf{x}, W) = W\mathbf{x} + \mathbf{b}$$

- Loss function
 - ▶ Quantifying *good* parameters
- Optimization
 - ▶ Finding optimal parameters to minimize the loss function
- Convolutional Neural Networks